Build Version 1.3.2 DEVELOPMENT: Technical Implementation Report

Project Name: AdjusterAssist™
Version: 1.3.2
Scope: Claim Drafting Scenario Improvement
Prepared By: Rudranarayan Sahu
Date: 29 April 2026

1. Introduction

Build 1.3.2 represents a significant architectural shift for AdjusterAssist™. This version moves the platform beyond simple response generation toward a persistent, intelligent conversational claim workspace. The primary objective of this build was to enhance the depth and professional accuracy of the drafting engine by introducing a multi-layered analysis pipeline and sophisticated context retention strategies. The system now behaves as a context-aware assistant capable of identifying risks, automating audience-specific formatting, and maintaining a high-fidelity memory of the entire claim lifecycle.

2. Implemented Features

A. Refined Payload Structure

The backend communication model has been optimized for efficiency and data integrity. This refined structure ensures that every request sent to the AI engine contains the necessary metadata to maintain professional guardrails.

- **Compliance Flag:** A dedicated parameter ensuring all generated drafts adhere to insurance industry regulations and internal carrier standards.
- **Attachment Context:** Improved handling of external data payloads (images and PDFs) to provide the AI with situational awareness of uploaded documentation.

B. Content Replacement Logic

A new UI-level logic has been integrated to handle refinements and variant generation. This feature allows for seamless updates to existing drafts, enabling users to iterate on specific sections of a response without losing the integrity of the original context. This is particularly useful during the creation of alternative variants for complex claims.

C. Input Analysis Layer

To ensure "premium" response quality, we have implemented a dedicated pre-processing layer that analyzes user input before drafting begins.

- **Facts Analysis:** Automatically extracts and validates key claim data points to ensure accuracy in the final output.

- **Risk Flag Trigger Logic:** A proactive safety system that scans inputs for potential inconsistencies or high-risk claim scenarios, alerting the engine to apply specific professional guardrails.

D. Audience Detection & Instruction Triggering

The system now features an intelligent audience classification engine that maps the output format to the intended recipient.

- **Automated Mapping:** The system identifies whether the audience is "Internal" or "External" based on the output classifier. For example, a 'File Note' triggers internal-facing instructions, while an 'Email' triggers external-facing templates.
- **Instruction Trigger Logic:** Based on the detected audience, the AI dynamically adjusts its tone, level of detail, and legal headers to match professional expectations.

E. Expanded Output Format Library

We have added three new professional output formats, complete with dedicated instruction sets and template packs:

Output Format	Description	Primary Audience
FNOL (First Notice of Loss)	Structured summaries for initial claim reporting.	Internal / Carrier
Attorney Correspondence	Formal communication templates tailored for legal representatives.	External
Inspection Summary	Concise reporting for field inspection findings and damage assessments.	Internal

F. Confidence Logic in Classifier

To ensure reliability, the Output Format Classifier now includes a confidence-scoring mechanism. If the AI is uncertain about the required format, it applies a secondary validation step or defaults to a safe, neutral template to prevent formatting errors.

G. RAG (Retrieval-Augmented Generation) Implementation

The core of the system's memory has been upgraded from raw history passing to a sophisticated RAG architecture for superior context management.

- **Ingestion Pipeline:** A continuous process that captures both user inputs and AI-generated responses, ensuring the entire claim history is indexed.
- **Chunks & Embeddings Rules:** Data is broken down into optimized semantic chunks and converted into vector embeddings, allowing the system to "remember" specific details based on meaning rather than just keywords.
- **Retrieval Pipeline:** During response generation, the system triggers a retrieval search to pull only the most relevant claim facts into the current prompt, minimizing "hallucinations" and maximizing precision.

3. Conclusion

The implementation of Build 1.3.2 has successfully transitioned AdjusterAssist™ into a robust, enterprise-ready drafting environment. By layering RAG technology with automated audience detection and proactive risk analysis, we have created a system that is both scalable and technically superior to standard LLM implementations. These improvements provide a solid foundation for the planned multi-role (Carrier vs. Public Adjuster) enhancements in future versions.

Please proceed with testing the updated features and provide feedback on the current build stability.